
A Description of Holland's Royal Road Function

Terry Jones

Santa Fe Institute
1399 Hyde Park Road
Santa Fe, NM 87501
terry@santafe.edu

Abstract

This paper contains a description of John Holland's Royal Road function, which was presented at the Fifth International Conference on Genetic Algorithms in July 1993 and posted to the Internet Genetic Algorithms mailing list in August 1993.

1. Introduction

The Royal Road functions were introduced by Mitchell, Forrest, and Holland (1992). They were designed as functions that would be simple for a genetic algorithm to optimize, but difficult for a hillclimber. However, a form of hillclimbing suggested by Palmer (personal communication, July 1992), easily outperformed a genetic algorithm on what had been expected to be the simplest Royal Road function for that algorithm. Recently, Holland (1993) presented a revised class of Royal Road functions designed to create insurmountable difficulties for a wider class of hillclimbers, and yet still be admissible to optimization by a genetic algorithm.

Holland used seven variables in the description he posted to the genetic algorithms community. In describing the class of functions, I will use his variable names where reasonable and indicate deviations. Holland suggested a default value for each of these variables; in what follows, I will use these values to provide illustration. A summary of these variables and their default values can be found in Section 3.

The main description here is of the function posted to the genetic algorithms community. Holland has also suggested further generalizations, which are discussed in Section 6.

The following terms will be used consistently in what follows: *block*, *complete*, *gap*, and *region*. These are introduced below and can be relied on to always mean exactly the same thing.

2. Description

The function takes a binary string as input and produces a real value. The function is used to define a search task in which one wants to locate strings that produce high function (fitness) values. Section 5 provides some more information on fitness values.

The string is composed of 2^k nonoverlapping contiguous regions, each of length $b + g$. With Holland's defaults, $k = 4$, $b = 8$, $g = 7$, there are 16 regions of length 15, giving an

Table 1. PART fitness values for Holland's default settings.

1s in block	0	1	2	3	4	5	6	7	8
Block fitness	0.00	0.02	0.04	0.06	0.08	-0.02	-0.04	-0.06	0.00

overall string length of 240. We will number the regions, from the left, as $0, 1, \dots, 2^k - 1$.

Each region is divided into two nonoverlapping pieces. The first, of length b , will be called the block. The second, of length g , will be called the gap.¹ In the fitness calculation, only the bits in the block part of each region are considered. The bits in the gap part of each region are completely ignored during fitness calculation. Holland's posting called the block part of each region by various names: "building blocks," "elementary building blocks," "schemata," "lower-level target schemata," and "elementary (lowest-level) building blocks (schemata)."

The fitness calculation proceeds in two steps. First, there is what Holland calls the PART calculation. Then follows the BONUS calculation. The overall fitness assigned to the string is the sum of these two calculations.

2.1 The PART Calculation

The PART calculation considers each block individually, and each in exactly the same fashion. Each block receives some fitness score, and these scores are added to produce the total PART contribution to the overall fitness.

The fitness of each block is based entirely on the number of 1 bits it contains.² The idea is to reward 1s up to a certain point. Every 1 up to (and including) a limit of m^* adds to the block's fitness by v . Thus, with the default settings, a block with three 1s would have fitness $3 * 0.02 = 0.06$. If a block contains more than m^* 1s, but less than b 1s, it receives $-v$ for each 1 over the limit. With the default settings $m^* = 4$, and so a block with six 1s is assigned a fitness of $(6 - 4) * -0.02 = -0.04$.

Finally, if a block consists entirely of 1s (i.e., it has b 1s), it receives *nothing* from the PART calculation. Such a block will be rewarded in the BONUS calculation. If a block consists entirely of 1 bits, it will be said to be *complete*.

From the above, we can construct a table of fitness values based on the number of 1s in a block. Table 1 gives the values for the default settings.

2.2 The BONUS Calculation

The idea of the BONUS calculation is to reward completed blocks and some combinations of completed blocks. Holland gives rewards for attaining what he calls "levels." At the lowest level,³ 0, rewards are given for complete blocks (i.e., blocks that consist entirely of 1s). If such a block exists, it receives u^* as its fitness. Any additional complete blocks receive a fitness of u . Thus, with the default settings, a string that contained three complete blocks would receive a BONUS fitness of $1.0 + (2 * 0.3) = 1.6$ for level 0.

But that is not the end of the BONUS story. At the next level, *pairs* of blocks are rewarded. Suppose we label the blocks from left to right as $B_0, B_1, \dots, B_{2^k-1}$. The function then rewards

1 Holland did not give a name for this part of the regions or give a variable name for its length. These gaps have been called "introns" (borrowing from biology) and have been used with varying degrees of success (Forrest & Mitchell, 1992; Levenick, 1991) to alter the effects of crossover in genetic algorithms.

2 Holland used $m(i)$ to denote the number of 1s in block i . I will not use a variable.

3 Holland numbered his lowest level 1, not 0.

Table 2. Holland's default variable settings.

Variable	Default
k	4
b	8
g	7
m^*	4
v	0.02
u^*	1.0
u	0.3

any completed pair (B_{2i}, B_{2i+1}) for $0 \leq i < 2^{k-1}$. The rewards are calculated in the same way as they were done at level 0. The first completed pair of blocks receives u^* , and additional completed pairs receive u .

Notice that not all pairs of completed blocks are so rewarded, only those whose indices are of the form $(2i, 2i + 1)$.

We can now calculate rewards for the next level, that of quadruplets of completed blocks, in the same fashion. In general, there are $k + 1$ levels, and at level $0 \leq l \leq k$, we are looking for contiguous sets of 2^l complete blocks, whose first block is labeled $B_{2^l i}$ with $0 \leq i < 2^{k-l}$. At all levels, the first such set of complete blocks receives fitness u^* , and additional sets of completed blocks receive fitness u . The total fitness for the level is the sum of these fitnesses.

To make this more concrete, consider the function with Holland's default values. With $k = 4$ we have 16 regions, and each contains a block of length $b = 8$. The BONUS fitness calculation rewards completed single blocks (this is level 0) and rewards the completion of the sets $\{B_0, B_1\}$, $\{B_2, B_3\}$, $\{B_4, B_5\}$, $\{B_6, B_7\}$, $\{B_8, B_9\}$, $\{B_{10}, B_{11}\}$, $\{B_{12}, B_{13}\}$, $\{B_{14}, B_{15}\}$, (level 1) $\{B_0, \dots, B_3\}$, $\{B_4, \dots, B_7\}$, $\{B_8, \dots, B_{11}\}$, $\{B_{12}, \dots, B_{15}\}$ (level 2) $\{B_0, \dots, B_7\}$, $\{B_8, \dots, B_{15}\}$ (level 3) and $\{B_0, \dots, B_{15}\}$ (level 4) of completed blocks.

The total BONUS contribution to the fitness is computed by adding the fitness at each of the $k + 1$ levels.

3. Summary of Variable Defaults and Meanings

Table 2 shows the default settings suggested by Holland. Here is a brief summary of each of these.

- k This determines the number of level 0 blocks. There will be 2^k of them. Equivalently, it determines the number of levels in the function, $k + 1$.
- b The number of bits in each block. These bits are examined when the fitness is calculated.
- g The number of bits in the gap following each block.⁴ These bits are ignored when the fitness is calculated.

⁴ The name is a slight misnomer, since the last block is followed by g bits which are not followed by another block.

- m^* This is used in calculating the PART fitness. It is the number of 1s that a block may have before being penalized for having too many. The penalty applies to blocks that have in excess of m^* 1s but not to those blocks that are complete.
- v This is also used in calculating the PART fitness. If the number of 1s in a block is m^* or less, then each receives v . If the number of 1s is greater than m^* but less than b , then each 1 over m^* *decreases* the fitness by v . If the block is complete, it receives no PART fitness.
- u^* The fitness acquired in the BONUS calculation by the first (if any) completed set of blocks at each level.
- u The fitness acquired in the BONUS calculation by the second and subsequent (if any) completed set of blocks at each level.

4. Relationships between Variables

The seven variables used to define this class of functions should be well understood before they are altered. They may be roughly divided into three groups:

- The variables k , b , and g can be thought of as structural variables. These can be changed at will; each adjusts a specific aspect of the function (see above); and they do not interact in any subtle way with any other variables.
- The PART variables are m^* and v .
- The BONUS variables are u^* and u .

The PART and BONUS variables *do* interact subtly, and various settings can radically alter the nature of the function thus defined. The function was originally supposed to reward up to a certain number of 1s in each block, then penalize any number greater than this, except for a complete block. If, for example, $u^* < vm^*$, this qualitative objective will not be satisfied. Here an optimal string (there are

$$2g2^k \binom{b}{m^*}^{2^k}$$

of them) will have m^* 1s in each block, since any more will result in a decreased fitness. Similarly, if $u < vm^*$ an optimal string will contain a single completed block and all others will have m^* 1s. This becomes even more odd when $u < vm^* < u^*$ since in this situation, the optimal bit pattern for a block depends on the contents of the other blocks—i.e., before any block is completed, the best a block can do is get completed; if there is exactly one completed block, the others will maximize the function by setting m^* bits; and finally, if there are several completed blocks, each of them would increase the string's fitness by moving to m^* 1s. But if they all do this, the pressure to return to a completed block returns!

These examples indicate that one should choose $vm^* < u \leq u^*$. There are a number of other more obvious inequalities among the variables that should be maintained. For example, $v \geq 0$ and $0 \leq m^* < b$.

5. Fitness Values with the Default Settings

With Holland's default settings, a maximally fit string has all 16 blocks complete. The gaps may contain any bit settings. Thus there are 2^{8^k} maximally fit strings. The fitness of these strings is 12.8. This can be calculated as follows. The PART fitness for each block is 0. The BONUS fitness at level 0 is

$$\begin{aligned} u^* + u(2^k - 1) &= 1.0 + 0.3 * 15 \\ &= 5.5 \end{aligned}$$

At level 1, we have eight completed pairs of blocks, for a total of $u^* + 7u = 3.1$; at level 2 we have four completed quadruples, giving $u^* + 3u = 1.9$; at level 3 we have two 8-tuples, for $u^* + u = 1.3$; and at level 4 we have $u^* = 1.0$. The total fitness is then

$$5.5 + 3.1 + 1.9 + 1.3 + 1.0 = 12.8$$

At the other end of the scale, it is possible for fitness to be negative. If every block contains $b - 1$ 1s, then each will receive $-v(b - 1 - m^*)$ in the PART calculation. The minimal fitness will therefore be the product of this value and 2^k . The BONUS calculation will not add anything since there are no complete blocks. With Holland's defaults, the minimum fitness is -0.9 .

6. Generalizations of the Function

There are several conceptually simple ways one can generalize the class of functions described above. Some of these are currently in use (Holland, personal communication, January 1994). This is by no means an exhaustive list; other possibilities, and some of the following, were presented by Forrest and Mitchell (1992). Once the basic function is understood, it is relatively simple to come up with plausible variations on the theme.

- The lengths of blocks and gaps need not be constant; one could specify 2^k different block and gap lengths.
- The maximal block bit pattern need not be all 1s. Holland suggests that each block's maximal pattern be randomly generated. To evaluate a string, one could initially XOR each block of the string with the maximal bit pattern for that block and then evaluate as normal.
- Holland has suggested the possibility of defining "potholes at every level" (Holland, personal communication, 1994). In this scheme, the BONUS fitnesses would be calculated in a manner akin to the PART fitnesses. At level 0, the PART calculation could proceed as usual. At every higher level, a number of components from the previous level would need to be completed. Suppose, for simplicity, that this number also happened to be b . Then fitness at each level above the lowest would be calculated using m^* and b (and possibly v or some other variable), as for the PART calculation at level 0. This would create a "pothole" at every level of the hierarchy.
- Blocks need not be contiguous. In the most general case, the block and gap bits could be spread randomly through the string. One would imagine that this would make the problem more difficult for a genetic algorithm using one- or two-point crossover (as

opposed to uniform crossover) since the building blocks would now not have short defining length.⁵

- The choice of which sets of blocks to reward in the BONUS calculation seems a little arbitrary. For example, there is no reward for getting any set of three blocks completed. Even at the levels that rewards are given, some sets are rewarded and others are not. For example, completing the set $\{B_0, B_1\}$ will increase a string's fitness at level 1, whereas completing the set $\{B_1, B_2\}$ will not. The choice of which sets to reward originates with the choices for R2 in Mitchell et al. (1992), which were made for simplicity.

7. Other Minor Points

Here are some relatively minor points.

- It is not possible to assign a setting to the variables to produce either of the functions R1 or R2 defined by Mitchell et al. (1992). In fact, the Royal Road functions, as defined in that paper, include the class of functions that Holland proposes. More than this, the class includes *all* functions $f : \{0, 1\}^n \rightarrow \mathbb{R}$ as one can assign any desired fitness to each of the 2^n schemata of order n to produce any desired function. But this is to ignore their intent, which is to provide a useful way to construct, think about, and concisely notate a class of functions of interest to people studying genetic algorithms. To these ends, they are most useful when used to express a function in which membership in a relatively small number of schemata is being rewarded.
- The intermediate building blocks of R2 are not present in R1 (Mitchell et al., 1992), making it possible to compare algorithm performance in the two situations. It is natural to ask how the hierarchical levels of Holland's class of functions are most naturally removed. One reasonable answer to this is to compute the fitness by doing the normal PART calculation and then doing only the lowest-level BONUS calculation.
- There doesn't seem to be any good reason for including the last gap on the string. If the purpose of such gaps is to introduce introns that (it is hoped) will alleviate hitchhiking problems, then the last gap could be omitted. Including the last gap does make the definition of the function simpler, though.
- Holland (1993) mentions "climbing levels" of the hierarchy and of "achieving" levels. These mean the same thing. A string that receives a BONUS fitness at level l (and not at any higher level) can be said to have reached or achieved level l . An algorithm may be said to be climbing to higher and higher levels if, over time, it finds strings that reach higher and higher levels. If one imagines the blocks as forming the base of a triangle, the pairs of blocks as forming the next level and so on up to the apex set (containing all blocks), then the progress of an algorithm, measured in terms of the highest level reached over time, can be thought of as a climb up these levels.

To say that an algorithm achieves a level therefore does not mean that it has found *all* the sets of blocks at a certain level. It means that it has found a single one. It is not possible to find all the sets of blocks at any level without simultaneously finding all sets

⁵ The expected defining length for a block of length b embedded randomly in a string of length n is $\mathcal{O}(n)$.

at all levels and thereby solving the entire problem. If the problem is not yet solved, there must exist at least one set of blocks at every level that is not complete.

8. Implementation

A C language implementation of Holland's function class and other information is available for anonymous ftp from [ftp.santafe.edu](ftp:santafe.edu) in the directory [pub/terry/jhrr.tar.gz](ftp://ftp.santafe.edu/pub/terry/jhrr.tar.gz).

Acknowledgments

Thanks to Derek Smith for useful conversations and for comments on the text, to Joseph Culberson and Ron Hightower for their implementations, and to Stephanie Forrest, John Holland, and Melanie Mitchell for answering questions about the precise nature of their functions.

References

- Forrest, S. & Mitchell, M. (1992). Towards a stronger building-blocks hypothesis: Effects of relative building-block fitness on GA performance. In *FOGA-92, Foundations of Genetic Algorithms* (pp. 109–126). Los Altos, CA: Morgan Kaufmann.
- Holland, J. H. (1993). Royal road functions. *Internet Genetic Algorithms Digest*, 7:22. Available in [ftp.santafe.edu:pub/terry/jhrr.tar.gz](ftp://ftp.santafe.edu/pub/terry/jhrr.tar.gz).
- Levenick, J. R. (1991). Inserting introns improves genetic algorithm success rate: Taking a cue from biology. In R. K. Belew and L. B. Booker (Eds.), *Proceedings of the Fourth International Conference on Genetic Algorithms* (pp. 123–127). San Mateo, CA: Morgan Kaufmann.
- Mitchell, M., Forrest, S., & Holland, J. H. (1992). The royal road for genetic algorithms: Fitness landscapes and GA performance. In F. J. Varela and P. Bourgine (Eds.), *Proceedings of the First European Conference on Artificial Life. Toward a Practice of Autonomous Systems* (pp. 245–254). Cambridge, MA: MIT Press.

This article has been cited by:

1. Álvaro Fialho, Luis Da Costa, Marc Schoenauer, Michèle Sebag. 2010. Analyzing bandit-based adaptive operator selection mechanisms. *Annals of Mathematics and Artificial Intelligence* . [\[CrossRef\]](#)
2. E.E. Korkmaz, G. Ucoluk. 2004. A Controlled Genetic Programming Approach for the Deceptive Domain. *IEEE Transactions on Systems, Man and Cybernetics, Part B (Cybernetics)* **34**:4, 1730-1742. [\[CrossRef\]](#)
3. Hideaki Suzuki, Yoh Iwasa. 1999. Crossover Accelerates Evolution in GAs with a Babel-like Fitness Landscape: Mathematical AnalysesCrossover Accelerates Evolution in GAs with a Babel-like Fitness Landscape: Mathematical Analyses. *Evolutionary Computation* **7**:3, 275-310. [\[Abstract\]](#) [\[PDF\]](#) [\[PDF Plus\]](#)